

Claude Code Is Getting More Powerful

Student Resource Guide: AI risk, three predictions, and the three skills that help you use powerful agents responsibly.

By: [Nate Herk](#)

1. Why Capability Changes the Stakes

The video opens with Anthropic CEO Dario Amodei describing his expectation that AI could exceed most or all human abilities. The practical concern is what happens when that intelligence can act through software, tools, and business systems.

Claude Code is an agent environment. An underlying model supplies intelligence, while connected tools and permissions determine what the agent can do. Reading a codebase, finding a flaw, and taking action create different risks from generating a text answer.

The Fable and Mythos shutdown

Three days after Anthropic released **Fable 5** and **Mythos 5**, a US directive restricted access by foreign nationals, including Anthropic employees inside the country. Anthropic said it could not reliably check nationality in real time, so it suspended both models for everyone. Source: [Anthropic access statement](#).

The reported trigger was a **jailbreak**, a technique that bypasses safeguards. It led Fable to identify software weaknesses and produce an exploit demonstration. Anthropic disputed the severity, describing minor, previously known vulnerabilities. Its later testing found that older Claude and GPT models could reproduce the demonstration.

Fable 5 and Mythos 5 share an underlying model. Fable has safeguards for general use; Mythos has fewer safeguards in selected areas and is restricted to trusted defensive partners. These model names should not be confused with the Claude Code application.

After 18 days, the restrictions were lifted. Anthropic reported a new **classifier** that blocked the specific technique in more than 99% of cases. It also committed to expanded government access before future releases and a shared framework for evaluating jailbreaks. Source: [Redeploying Fable 5](#).

Scope matters: Blocking one reported technique in more than 99% of cases does not mean an agent is 99% safe across every task or attack.

Vulnerabilities that existing checks missed

Anthropic reported that **Mythos Preview** had found thousands of previously unknown vulnerabilities and developed many related exploits, often autonomously. Two examples from the video illustrate the capability:

- OpenBSD: A vulnerability that had existed for 27 years.
- FFmpeg: A vulnerability that had existed for 16 years in code automated tests had reached five million times.

Source: [Anthropic research on Mythos Preview](#). The same ability can help defenders find problems and help attackers exploit them. Which tools and systems an agent can access becomes a central business decision.

Why one shutdown does not stop the race

The video places the shutdown alongside earlier government action against Anthropic and continued progress from Chinese labs and other American companies. Nate's interpretation is that competition makes unilateral slowing difficult.

Evidence boundary: The video does not claim public evidence that competition caused the government to lift the restrictions. The race argument is an interpretation of the wider incentives.

2. Three Predictions to Watch

These are Nate's predictions about where AI adoption may lead. Existing programs and employment projections provide context, but the future outcomes are not guarantees.

Prediction 1: AI literacy becomes workforce infrastructure

AI literacy may become a basic workforce skill over the next 5 to 10 years. Governments have incentives to help workers and businesses use AI as other countries invest in the same capabilities.

The White House has established an AI education task force. NSF's AI Ready America initiative supports literacy, workforce skills, and practical adoption by businesses and communities. Sources: [White House AI education order](#); [NSF AI Ready America](#).

- Workers may receive publicly funded AI training through colleges or workforce programs.
- Teachers may receive AI training through normal professional development.
- Local programs may help small businesses adopt AI.
- Schools, libraries, and community colleges may become regular access points for AI education.

Prediction 2: Cybersecurity grows, and the work changes

The Bureau of Labor Statistics projects 29% employment growth for information security analysts from 2024 to 2034. This is an occupation forecast, not a guarantee about any individual job. Source: [BLS Occupational Outlook Handbook](#).

Nate expects AI to handle more initial alert review and vulnerability scanning. Human work increasingly centers on verifying findings, controlling access, investigating unusual cases, responding to incidents, and demonstrating that systems are secure.

Cyber insurance may make agent oversight a routine business requirement. The video highlights existing security fundamentals: **multifactor authentication**, software updates, backups, employee training, vendor access, and incident response plans.

Nate predicts that insurers will increasingly ask which agents can move money, access customer data, or deploy code; whether their actions are logged; and who can shut them down. Poor answers could make coverage more costly or harder to obtain.

Insurance is a financial backstop: Coverage may help with losses and response costs, depending on the policy. It does not prevent attacks or replace security controls.

Controlling your own agents also does not stop outside attackers from using AI. The video points to more convincing phishing, faster weakness discovery, and faster movement inside compromised systems.

Prediction 3: AI arms control emerges unevenly

Nate expects attempts at international controls, but doubts that every country and lab will agree to a clean global pause. The video points to international AI agreements and UN work on verifying claims about advanced AI.

Training a leading model from scratch requires specialized chips, large computing clusters, power, infrastructure, and substantial funding. Using or adapting an existing downloadable model requires far fewer resources. Once model copies spread, control becomes harder.

- Possible oversight tools include chip tracking, cloud logs, and model evaluations.
- Incident reports and data center inspections could help make major training runs visible.
- These controls cannot prove that every lab, country, or private actor has stopped.

Governments may seek stronger safety rules while also funding domestic computing, training workers, and reducing dependence on foreign models. Those competing incentives make coordination difficult.

How to interpret the warnings in the video

The transcript discusses Jacob Coxon's resignation after research work at OpenAI and Anthropic. He argued that the labs were racing toward systems that could improve themselves without acting

responsibly. The video also shows Bernie Sanders announcing an intention to introduce legislation to ban superintelligence and pause development.

It quotes Evan Hubinger's personal estimate of a greater than 10% chance of human extinction from AI within a decade and his concern that alignment for superintelligence remains unsolved. These are attributed warnings and personal forecasts, not measured probabilities or proof that an announced bill became law.

3. Skill One: Find the Pain

Core principle: Understand which business problem deserves automation before choosing or building an agent.

A client may request an agent they saw online. Building exactly what they ask for can automate the wrong process. More capable models make good diagnosis more valuable because they can execute a mistaken plan quickly.

The discovery question

Ask: If your business grew 10 times tomorrow, which process would break first?

Use the answer to investigate the **constraint**, the part of the business that is limiting results. Separate the desired outcome, such as more revenue, from the client's proposed solution, such as more leads.

The med spa example

- Initial request: Generate more leads to increase revenue.
- Actual situation: Plenty of leads were already booking appointments.
- Constraint: Too many booked customers failed to show up.
- Better first project: Improve follow up and reduce missed appointments.

Adding leads would have sent more people into a process that was already losing revenue. Addressing attendance targeted the existing loss. The transcript reports no numerical before and after results, so none are assumed here.

Apply it to your next project

1. State the business outcome the stakeholder wants.
2. Ask which process would break if demand grew 10 times.
3. Check where the current process loses customers, time, or money.
4. Choose the first automation that directly addresses that constraint.

4. Skill Two: Control the System

Agent access is not an all or nothing decision. Connecting email, a CRM, payment tools, and a codebase without planning permissions can give a new agent more authority than you would give a new employee.

Core principle: More capable agents need tighter boundaries.

Define the boundaries before connecting tools

- Read access: Which records, messages, and files may the agent see?
- Write access: Which fields, files, or systems may it change?
- Approval: Which actions must wait for a human decision?
- Traceability: Where can you review exactly what the agent did?
- Shutdown: Who can stop the agent quickly, and how?

Use the operating rules from the video

1. Begin with read only access wherever possible.
2. Have the agent prepare drafts before sending messages or changing live systems.
3. Require human approval for money, deletions, publishing, and production changes.
4. Log actions so you can reconstruct what happened.
5. Make shutdown practical and assign someone to operate it.

You should be able to explain which agents touch customer data, who approves sensitive actions, where the logs are stored, and what happens if an agent is compromised. These controls address risks from your own agents while supporting the broader security program.

5. Skill Three: Own the Outcome

An agent can produce a polished answer that is wrong. It can quote the wrong price, email the wrong person, or make a bad recommendation. A person must remain accountable for checking its work and handling the consequences.

Define success before the build

Choose the business number the system is supposed to improve. The video suggests acquiring more customers, increasing the value of each customer, reducing costs, or preventing expensive mistakes.

- **Baseline:** The current value of the chosen metric.
- **Target:** The result you want to reach.
- **Timeframe:** How long the system has to reach the target.
- **Failure review:** What could go wrong, and who is responsible for checking?

Measure results and review behavior

For the med spa, record the appointment attendance rate before introducing the follow up system. Set a target and measure the same metric after launch. Use a consistent definition so the comparison is meaningful.

Because the agent communicates with customers, review what it says, define when a person takes over, and retain a log of every message and interaction. A better business metric does not by itself demonstrate that the agent behaved appropriately.

Two questions to answer: Did the business number improve? Can you show how the system behaved, who approved its actions, and what happened when it was unsure?

Owning the outcome includes taking responsibility for mistakes and making successful results visible. Your value is in selecting useful work, controlling execution, and standing behind the result.

6. Your Next Project Worksheet

Complete this short worksheet before building or expanding an agent. It applies the three skills taught in the video.

Find the pain

- Desired business outcome: _____
- Process that would break at 10 times the volume: _____
- Evidence of the current constraint: _____
- First workflow to improve: _____

Control the system

- Allowed data and read access: _____
- Allowed changes and draft outputs: _____
- Actions needing approval and the approver: _____
- Log location, shutdown method, and responsible person: _____

Own the outcome

- Metric and baseline: _____
 - Target and review timeframe: _____
 - Likely mistakes and the person checking them: _____
 - When a human takes over: _____
 - Evidence to show the stakeholder: _____
-

Key Takeaways

1. Powerful models become more consequential when agents can access tools, data, and live systems.
 2. A restriction on one model does not stop competitors or remove the underlying incentives to advance.
 3. Treat expanded AI education, changing security careers, and AI arms control as predictions to watch.
 4. Find the actual business constraint before deciding what to automate.
 5. Start with limited access, use drafts and approvals, log actions, and maintain a shutdown path.
 6. Establish a baseline, target, timeframe, and human owner before launch.
 7. Demonstrate both the business result and how the agent behaved. Responsibility stays with people.
-

Want to connect with others building and monetizing AI automation?

[Become an AIS Plus Member](#)